

An In-Depth Look of BFT Consensus in Blockchain: Challenges and Opportunities

(System)

Suyash Gupta,
Sajjad Rahnama,

Jelle Hellings,
Mohammad Sadoghi




Agenda

Session I

- 1) Blockchain 101
 - 1) What is Blockchain, Applications and Components?
 - 2) Permissionless and Permissioned Blockchain.
- 2) Transactions and Consensus
- 3) Primer on Byzantine Fault-Tolerant Consensus
- 4) Existing Optimizations for BFT Consensus.

Session II


- 1) PoE: Two-Phases Resilient Consensus.
- 2) MultiBFT: Parallel and Wait-free Consensus
- 3) GeoBFT: Global Scale Consensus
- 4) Reducing Communication between Clusters




Hands-on ResilientDB

What is Blockchain?

- A linked list of blocks.
- Each block contains hash of the previous block.
- A block contains information about some client transactions.




Why Blockchain?



- By User:Pedant, User:Wapcaplet, User:Antonu, User:Vanderlindenma, User:.js. - Composition of File:Barnstar of Diligence Hires.png + File:Voting hand.svg., CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=45960536>

- <https://blog.devolutions.net/2017/10/whats-the-difference-between-2fa-and-mfa>

Blockchain Applications?




Components of a Blockchain System

- Replicas → Store all the data.
- Client → Sends transactions to process.
- Consensus Protocol → Helps ordering transactions.
- Cryptographic Constructs → Authenticate replicas and clients.
- Ledger → Records transactions.



Consensus



Types of Blockchain Systems

- **Permissionless** → **Open Access**
 - Anyone can participate.
 - Identities of the replicas unknown.
 - Applications include crypto-currency and money exchange.
- **Permissioned** → **Restricted Access**
 - Only a selected group of replicas, although untrusted can participate.
 - Identities of the replica known a priori.
 - Applications include health-care and energy trading.


BITCOIN

- First Crypto-currency → a monetary application.
- Uses Nakamoto consensus → Proof-of-Work beneath the skin.
- Supports permissionless access.
- Requires solving hard cryptographic puzzles.
- Any replica that wants to create a new block proves that it did solve the puzzle.
- Difficulty of the puzzle helps prevent malicious attacks.



ETHEREUM

- Another Crypto-currency → a token used in variety of applications.
- Uses Proof-of-Work but plans to start using Proof-of-Stake.
- Supports permissionless access.
- Allows programmers to design their transactions or “*smart contracts*”.
- Hard dependency on Ethereum Virtual machine (EVM).
- Envisions design of Permissioned applications.



Terrorists Turn to Bitcoin for Funding, and They're Learning Fast



Login

Search Q

Disrupt Berlin 2019

THE VERGE

TECH ▾ REVIEWS ▾ SCIENCE ▾ CREATORS ▾ ENTERTAINMENT ▾ VIDEO FEATU

REPORT \ TECH \ CYBERSECURITY


Why the Ethereum Classic hack is a bad omen for the blockchain

The 51 percent attack is real, and it's easier than ever

By Russell Brandom | Jan 9, 2019, 8:47am EST



News ▾ Features ▾ Price Analysis ▾ Market Tools ▾ Cryptopedia ▾ Industry ▾



By William Suberg

NOV 16, 2019

Bitcoin Cash Hard Fork Sees Miners 'Waste' Money on 14 Invalid Blocks



Bitcoin 24h
\$7,355.40 -2.58%

Ethereum 24h
\$151.13 -0.62%

XRP 24h
\$0.223687 -


Story from Tech →

Bitcoin Cash Miners Undo Attacker's Transactions With '51% Attack'

May 24, 2019 at 21:17 UTC • Updated May 25, 2019 at 10:39 UTC



Permissioned Blockchain Systems

- Require identities of the participating replicas to be known a priori.
- Replicas still untrusted → Consensus through traditional BFT protocols.
- Computationally in-expensive.
- Communication intensive.
- Prevent chain forks.
- Suitable for needs of an industry → JP Morgan, IBM, Oracle
- Advent path for *Blockchain Databases*.





Transactions, Agreement and Consensus




The Omniscient Transaction

- A transformation from a *consistent* state to another consistent state.
- A *contract* between two or more parties.
- A collection of *Read* or *Write* operations.
- Types of transactions: nested, compensating, multi-operation etc.

ACID Properties

- **Atomicity:** A transaction either completes fully or none of its changes take place.
- **Consistency:** The transaction must obey legal protocols
- **Isolation:** The intermediate state of a transaction is invisible to other transactions
- Durability: Once a transaction is committed, it cannot be abrogated





Consistency vs Availability

- An ongoing struggle that causes *performance tradeoffs*.
- Availability → Database needs to be always available for use.
 - Solution? Replication
 - Issues? Faults, Failures and Attacks.
- Consistency → Database needs to be correct.
 - Solution? All replicas should have same state.
 - Issues? Expensive.



Partitioning vs Replicating

- Distributed Databases can be partitioned, replicated or both.
- Partitioning → Split database into multiple disjoint partitions.
- Replication → Multiple full copies of the database.
- Partitioned Replication → Multiple partitions, where each partition employs replication.



Agreement in Partitioned Databases

- Partitioned Databases receive client transactions that may access multiple partitions.
- Deciding the fate of *multi-partition transactions* requires coordination among the partitions.
- *Coordination* is costly but necessary.
- Coordination or agreement among the partitions should be both *safe* and *live*.

A Deep Dive into BFT Consensus (Theory Slides Continue)



UCDAVIS
UNIVERSITY OF CALIFORNIA

Modern BFT Consensus Optimizations




UCDAVIS
UNIVERSITY OF CALIFORNIA


PBFT: Practical Byzantine Fault Tolerance

- First practical Byzantine Fault Tolerant Protocol.
- Tolerates up to f failure out of $3f+1$ replicas
- Three phases of which two require quadratic communication complexity.
- Safety is always guaranteed and Liveness is guaranteed in periods of partial synchrony.
- View-Change protocol for replacing malicious primary

PBFT Failure-Free Flow




PBFT Primary Failure (View Change)



All about Eve: Execute-Verify Replication for Multi-Core Servers [OSDI'12]

- New Architecture: **Execute/verify** instead of Agree/Execute.
- Execute multiple requests concurrently and then verify the output.
- Takes advantage of **parallel hardware** to improve performance.
- Non-deterministic multi-threaded execution
- The **Byzantine agreement is on the output** instead of sequence.
- Allows **divergence** in execute and roll back in case

Eve's Execute-Verify Flow




Zyzyva: Speculative Byzantine Fault Tolerance


[SOSP'07]

- Employs **Speculation** to achieve consensus in a **single** phase.
- In the best case (no failures), it only requires **linear** communication complexity.
- **Depends** on its good clients, for achieving common order among the replicas.
- Client **needs** identical response from all the **$3f+1$** replicas.
- With just **one crash fault** it faces severe throughput degradation.
- Recently, proven **unsafe**.


Zyzzya Failure-Free Flow



Zyzzya Flow with the Failure of One Non-Primary Replica



Zyzzzyva View Change Protocol




RBFT: Redundant Byzantine Fault Tolerance

[ICDCS'13]


- Robust BFT protocol, perform well in the presence of **smart malicious primary**.
- Runs **f+1 instance of BFT** protocol to monitor best performance .
- Do not rely on **one specific primary**.
- One **Master primary** and f backup instance of protocol always being executed.
- **Goal:** replicas monitor the throughput of the primary and replace it with one of the backups when it is slow to achieve robustness

RBFT's Multiple Redundant Primary Design



$n=7$ replicas, $f=1$ and $f+1=3$ instances

RBFT Failure-Free Flow




Hotstuff: BFT Consensus in the Lens of Blockchain

[PODC'19]

- **Splits** each $O(n^2)$ phase of PBFT into two linear phases.
- **Changes** leader at the end of each consensus.
- Employs **expensive** threshold signatures to linearize consensus.
- **Two versions:**
 - **Basic Hotstuff:** Unfolding each phase of PBFT into **two round** and make it linear.
 - **Chained Hotstuff:** Pipelined version of basic one, each phase different role for different view.

Hotstuff Protocol Flow



Attested Append-Only Memory: Making Adversaries Stick to their Word [SOSP'07]

- Uses a trusted component to reduce the hard-limit of one-third byzantine failures.
- Trusted component **removes equivocation** → Primary **cannot lie** about the order.
- A2M → Set of trusted, undeniable, ordered logs.
- Messages in the log can be verified by everyone using attestation.
- **Attestation overheads:** log writing, verifying.

Attested Append-Only Memory


A2M-PBFT-E

- Protecting the execution by adding attestation to client reply
- Safety and liveness when $f < n/3$
- **Just Safety** when $f < n/2$

A2M-PBFT-EA


- Append all messages to the A2M Log
- Protect Execution and Agreement with log
- Safety and liveness when $f < n/2$

Attested Append-Only Memory: PBFT with Execution Protection



Attested Append-Only Memory: Making Adversaries Stick to their Word

Attested Append Only Memory PBFT with Execution and Agreement Protection:




Towards Scaling Blockchain Systems via Sharding

[SIGMOD'19]

- Introduces the notion of **multiple chains**.
- Data is partitioned into shards → Each shard uses $3f+1$ replication.
- PBFT within each shard to provide Byzantine Fault-Tolerance.
- Multi-shard transactions require **Two-Phase Commit** protocol.
- Authors use SGX Trusted Hardware to reduce costs ($n = 2f+1$)


Towards Scaling Blockchain Systems via Sharding



Single-shard transactions are fast!

Towards Scaling Blockchain Systems via Sharding


Multi-shard Transactions need 2PC protocol → Initiated by the Reference Committee.



Monoxide: Scale out Blockchains with Asynchronous Consensus Zones (NSDI'19)

- Apply sharding on public blockchain cryptocurrency with Asynchronous Consensus Zone.
- Eventual Atomicity: First withdraw transaction, Later the deposit transaction.
- Asynchronous Consensus Zone: Parallel chains in different zones working independently.
- Mining power amplification: Chu-ko-nu Mining distributes mining power evenly across zones.
- Cross Zone Transactions: broadcasting relay transactions for the inter-zone part of the TXNs.


Monoxide: Scale out Blockchains with Asynchronous Consensus Zones




Normal Case with Good Primary and one fault

Algorand: Scaling Byzantine Agreement for Cryptocurrencies [SOSP'17]

- **Committee Based Consensus:** Scalability through consensus among selected users.
- **Proof of Stake:** Block proposers selected based on their stake.
- **Cryptographic Sortition:** Committee selection done independently at each node.
- **Verifiable Random Function:** Takes a secret key and a value and produces a pseudorandom output, with a proof



Algorand: Scaling Byzantine Agreement for Cryptocurrencies



SESSION II

An In-Depth Look of BFT Consensus in Blockchain: Challenges and Opportunities



UCDAVIS
UNIVERSITY OF CALIFORNIA



Requirements of Existing BFT Protocols

- 1) Require three phases of communication, of which two necessitate quadratic communication (PBFT).
- 2) Expect no failures or dependence on clients (Zyzzyva).
- 3) Incur high client latencies due to many phases of communication (PBFT, HotStuff).
- 4) Require threshold signatures, which are computationally expensive (HotStuff).
- 5) Require more than $3f+1$ replicas (Q/U, HQ).
- 6) Need trusted components (AHL, Attested Append-only memory).


Proof-of-Execution (PoE): Reaching Consensus through Fault-Tolerant Speculation

- *Speculative Execution* to reduce the client latency.
- *Out-of-Order message processing* for transactions.
- *Two-Phases* of communication.
- *No Dependence* on Clients or requirement of expensive cryptographic primitives.


PoE vs Other Protocols

Protocol	Phases	Messages	Computation	Resilience	Requirements
ZYZZYVA	1	$O(n)$	high	0	reliable clients
PoE (our paper)	2	$O(n + n^2)$	low	f	
PBFT	3	$O(n + 2n^2)$	low	f	
HOTSTUFF	4	$O(n + 3n^2)$	high	f	
HOTSTUFF-TS	8	$O(4n)$	very high	f	threshold sign.


PoE Protocol (Non-Faulty Primary)



PoE View Change Protocol: Replacing Malicious Primary



PoE Scalability under Single Failure




Scaling Blockchain Databases through Parallel Resilient Consensus Paradigm


- Why should BFT protocols rely on just *one* primary replica?
- Malicious primary can *throttle* the system throughput.
- Malicious primary requires *replacement* —> fall in throughput.

Multiple Byzantine Fault-Tolerance (MultiBFT) Paradigm

- Designate multiple replicas as Primaries!
- Run multiple parallel consensuses on each replica.





MultiBFT Flow



Malicious Primaries Collusion


- Multiple malicious primaries can prevent liveness!
- Solution → Optimistic Recovery through State Exchange.





MultiBFT using PBFT with 2 parallel instances on each replica


MultiBFT Scalability



Global Scale Resilient Blockchain Fabric


- Traditional BFT protocols do not scale to geographically large distances.
- Blockchain requires decentralization → replicas can be far apart → expensive communication!
- The underlying BFT consensus protocol should be topology-aware.


Vision Geo-Scale Byzantine Fault-Tolerance



GeoBFT Protocol

- GeoBFT groups replicas into clusters based on the distance between these replicas.
- Each cluster runs the PBFT consensus protocol, in parallel and independently.






GeoBFT Takeaways

- To ensure common ordering → linear communication among the clusters is required.
- Primary replica at each cluster sends a secure certificate to $f+1$ replicas of every other cluster.
- Certificates guarantee common order for execution.
- If primary sends invalid certificates → will be detected as malicious.

GeoBFT Scalability



ResilientDB: High Throughput Yielding, Scalable Permissioned Blockchain Fabric

Visit at: <https://resilientdb.com/>




Why Should You Choose ResilientDB?

- 1) Bitcoin and Ethereum offer low throughputs of *10 txns/s*.
- 2) Existing Permissioned Blockchain Databases still have low throughputs (*20K txns/s*).
- 3) Prior works blame BFT consensus as *expensive*.
- 4) System Design is mostly *overlooked*.
- 5) ResilientDB adopts *well-researched* database and system practices.


Dissecting Existing Permissioned Blockchains

- 1) Single-threaded Monolithic Design
- 2) Successive Phases of Consensus
- 3) Integrated Ordering and Execution
- 4) Strict Ordering
- 5) Off-Chain Memory Management
- 6) Expensive Cryptographic Practices


Can a well-crafted system based on a classical BFT protocol outperform a modern protocol?




ResilientDB Architecture



ResilientDB Multi-Threaded Deep Pipeline




Pipeline at Primary Replica




Pipeline at Non-Primary Replica

Insight 1: Multi-Threaded pipeline Gains




Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

Insight 1: Multi-Threaded pipeline Gains




Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

Insight 1: Multi-Threaded pipeline Gains





Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

Insight 1: Multi-Threaded pipeline Gains





Parallelizing and Pipelining tasks across worker, execution (E) and batch-threads (B).

Insight 2: Optimal Batching Gains



More transactions batched together → increase in throughput
→ reduced phases of consensus.

Insight 3: Memory Storage Gains



In-memory blockchain storage → reduces access cost.

ResilientDB: Hands On

Visit at: <https://github.com/resilientdb/resilientdb>



How to Run ResilientDB?

- Go to <https://github.com/resilientdb/resilientdb> and Fork it!
- Install Docker-CE and Docker-Compose (Links on git)
- Use the Script "*resilientDB-docker*" as following:

./resilientDB-docker --clients=1 --replicas=4

./resilientDB-docker -d [default 4 replicas and 1 client]

- Result will be printed on STDOUT and stored in *res.out* file.

How to Run ResilientDB?

resilientdb / resilientdb

Watch 5 Unstar 11 Fork 13

Code Issues 1 Pull requests 0 Actions Projects 0 Wiki Security Insights

ResilientDB: A scalable permissioned blockchain fabric

46 commits 1 branch 0 packages 2 releases 4 contributors MIT

Branch: master New pull request Create new file Upload files Find file Clone or download

gupta-suyash readme updated Latest commit f2302e6 3 days ago

benchmarks	Initial Commit	16 days ago
blockchain	ledger archiecture defined	4 days ago
client	Initial Commit	16 days ago
deps	Initial Commit	16 days ago
scripts	added -e to handle multiple clients in docker-ifconfig	13 days ago
statistics	Initial Commit	16 days ago
system	ledger archiecture defined	4 days ago
transport	Initial Commit	16 days ago
.gitignore	Initial Commit	16 days ago
CHANGELOG.md	changelog added	3 days ago
CODE_OF_CONDUCT.md	Create CODE_OF_CONDUCT.md	15 days ago
LICENSE.md	Initial Commit	16 days ago
Makefile	Initial Commit	16 days ago
README.md	readme updated	3 days ago
config.cpp	Initial Commit	16 days ago
config.h	ledger archiecture defined	4 days ago
resilientDB-docker	Initial Commit	16 days ago

Docker CE

What is Docker?

*an open-source project that automates the deployment of software applications inside **containers** by providing an additional layer of abstraction and automation of OS-level virtualization on Linux.*

- Run a distributed program on one machine
- Simulate with lightweight virtual machines

Docker CE

What is Docker?

*an open-source project that automates the deployment of software applications inside **containers** by providing an additional layer of abstraction and automation of OS-level virtualization on Linux.*

- Run a distributed program on one machine
- Simulate with lightweight virtual machines


Resilient DB

```
./resilientDB-docker -d
```

- Remove old Containers
- Create new Containers
- Create IP address settings
- Install dependencies
- Compile Code
- Run binary files
- Gather the results

```
sajjad@sajjad-xps:~/WS/expo/resilientdb|master⚡
> ./resilientDB-docker -d
Number of Replicas: 4
Number of Clients: 1
Stopping previous containers...
Stopping s3 ... done
Stopping s1 ... done
Stopping s4 ... done
Stopping c1 ... done
Stopping s2 ... done
Removing s3 ... done
Removing s1 ... done
Removing s4 ... done
Removing c1 ... done
Removing s2 ... done
Removing network resilientdb_default
Successfully stopped
Creating docker compose file ...
Docker compose file created --> docker-compose.yml  Send problem report to the developers?
Starting the containers...
Creating network "resilientdb_default" with the default driver
Creating s4 ... done
Creating c1 ... done
Creating s1 ... done
Creating s2 ... done
Creating s3 ... done
ifconfig file exists... Deleting File
Deleted
Server sequence --> IP
c1 --> 172.21.0.3
s1 --> 172.21.0.4
s2 --> 172.21.0.6
s3 --> 172.21.0.2
s4 --> 172.21.0.5
Put Client IP at the bottom
ifconfig.txt Created!

Checking Dependencies...
Installing dependencies..
/home/sajjad/WS/expo/resilientdb
Dependencies has been installed
```




Resilient DB

- Throughput
 - Transaction per second
- Average Latency
 - The from client request to client reply
- Working Thread idleness
 - The time that thread is waiting
- WT0: Consensus Messages
- WT1 and WT2: Batch Threads
- WT3: checkpointing Thread
- WT4: Execute Theread

```
Throughputs:  
0: 38525  
1: 38530  
2: 38558  
3: 38551  
4: 38564  
Latencies:  
latency 4: 0.505870  
  
idle times:  
Idleness of node: 0  
Worker THD 0: 116.227  
Worker THD 1: 62.0772  
Worker THD 2: 62.2130  
Worker THD 3: 105.098  
Worker THD 4: 74.9193  
Idleness of node: 1  
Worker THD 0: 39.3157  
Worker THD 1: 0.00000  
Worker THD 2: 0.00000  
Worker THD 3: 104.700  
Worker THD 4: 74.8603  
Idleness of node: 2  
Worker THD 0: 35.0847  
Worker THD 1: 0.00000  
Worker THD 2: 0.00000  
Worker THD 3: 102.415  
Worker THD 4: 78.1078  
Idleness of node: 3  
Worker THD 0: 38.4452  
Worker THD 1: 0.00000  
Worker THD 2: 0.00000  
Worker THD 3: 107.512  
Worker THD 4: 77.6965  
Memory:  
0: 172 MB  
1: 156 MB  
2: 155 MB  
3: 156 MB  
4: 812 MB  
  
avg thp: 4: 38541  
avg lt : 1: .505  
Code Ran successfully ---> res.out
```

PBFT: Practical Byzantine Fault Tolerance



PBFT: Practical Byzantine Fault Tolerance

Client Request

- Client/client_main.cpp
- System/client_thread.cpp
- ClientQueryBatch Class
- Process ClientBatch in primary

```
C++ client_main.cpp ×  
client > C++ client_main.cpp > ...  
31 int main(int argc, char *argv[]){  
32     printf("Running client...\n\n");  
33     // 0. initialize global data structure  
34     parser(argc, argv);  
35     assert(g_node_id >= g_node_cnt);  
36     uint64_t seed = get_sys_clock();  
37     srand(seed);  
38     printf("Random seed: %ld\n", seed);  
39  
40     int64_t starttime;  
41     int64_t endtime;  
42     starttime = get_server_clock();  
43     // per-partition malloc  
44     printf("Initializing stats... ");  
45     fflush(stdout);  
46     stats.init(g_total_client_thread_cnt);  
47     printf("Done\n");  
48     printf("Initializing transport manager... ");  
49     fflush(stdout);  
50     tport_man.init();  
51     printf("Done\n");  
52     printf("Initializing client manager... ");  
53     Workload *m_wl = new YCSBWorkload;  
54     m_wl->Workload::init();  
55 }
```

```
C++ client_thread.cpp ×  
system > C++ client_thread.cpp > ...  
79  
80 RC ClientThread::run()  
81 {  
82     tsetup();  
83     printf("Running ClientThread %d\n", _thd_id);  
84  
85     while (true)  
86     {  
87         keyMTX.lock();  
88         if (keyAvail)  
89         {  
90             keyMTX.unlock();  
91             break;  
92         }  
93         keyMTX.unlock();  
94     }  
95  
96     BaseQuery *m_query;  
97     uint64_t iters = 0;  
98     uint32_t num_txns_sent = 0;  
99     int txns_sent[g_node_cnt];  
100    for (uint32_t i = 0; i < g_node_cnt; ++i)  
101        txns_sent[i] = 0;  
102  
103    run_starttime = get_sys_clock();  
104 }
```

PBFT: Practical Byzantine Fault Tolerance

Process Messages

- Transport/message.cpp
- System/worker_thread.cpp
- System/worker_thread_pbft.cpp
- Worker Thread: Run function
- Worker Thread: Process function

```
C++ worker_thread.cpp ×
system > C++ worker_thread.cpp > WorkerThread::run()

626 /**
627  * Starting point for each worker thread.
628 *
629 * Each worker-thread created in the main() starts here. Each worker-thread is alive
630 * till the time simulation is not done, and continuously perform a set of tasks.
631 * These tasks involve, dequeuing a message from its queue and then processing it
632 * through call to the relevant function.
633 */
634 RC WorkerThread::run()
635 {
636     tsetup();
637     printf("Running WorkerThread %d\n", _thd_id);
638
639     uint64_t agCount = 0, readystarttime, idlestarttime = 0;
640
641     // Setting batch (only relevant for batching threads).
642     next_set = 0;
643
644     while (!simulation->is_done())
645     {
646         txn_man = NULL;
647         heartbeat();
648         progress_stats();
649
650 #if VIEW_CHANGES
651         // Thread 0 continuously monitors the timer for each batch.
652         if (get_thd_id() == 0)
653         {
654             check_for_timeout();
655         }
656
657         if (g_node_id != get_current_view(get_thd_id()))
658         {
659             check_switch_view();
660         }
661     #endif
662
663         // Dequeue a message from its work_queue.
664         Message *msg = work_queue.dequeue(get_thd_id());
```

```
C++ worker_thread.cpp ×
system > C++ worker_thread.cpp > WorkerThread::process(Message *)
87 void WorkerThread::process(Message *msg)
88 {
89     RC rc __attribute__((unused));
90
91     switch (msg->get_rtype())
92     {
93     case KEYEX:
94         rc = process_key_exchange(msg);
95         break;
96     case CL_BATCH:
97         rc = process_client_batch(msg);
98         break;
99     case BATCH_REQ:
100        rc = process_batch(msg);
101        break;
102    case PBFT_CHKPT_MSG:
103        rc = process_pbft_chkpt_msg(msg);
104        break;
105    case EXECUTE_MSG:
106        rc = process_execute_msg(msg);
107        break;
108 #if VIEW_CHANGES
109    case VIEW_CHANGE:
110        rc = process_view_change_msg(msg);
111        break;
112    case NEW_VIEW:
113        rc = process_new_view_msg(msg);
114        break;
115 #endif
116    case PBFT_PREP_MSG:
117        rc = process_pbft_prep_msg(msg);
118        break;
119    case PBFT_COMMIT_MSG:
120        rc = process_pbft_commit_msg(msg);
121        break;
122    default:
123        printf("Msg: %d\n", msg->get_rtype());
124        fflush(stdout);
125        assert(false);
126    }
127 }
```

PBFT: Practical Byzantine Fault Tolerance

Process Client Message

- System/worker_thread_pbft.cpp
- process_client_batch Function
- Create and Send Batch Request
 - create_and_send_batchreq Function
 - Create Transactions
 - Create Digest
- BatchRequest Class
 - Pre-Prepare Message

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > ...
18  /**
19   * Processes an incoming client batch and sends a Pre-prepare message to all replicas.
20   *
21   * This function assumes that a client sends a batch of transactions and
22   * for each transaction in the batch, a separate transaction manager is created.
23   * Next, this batch is forwarded to all the replicas as a BatchRequests Message
24   * which corresponds to the Pre-Prepare stage in the PBFT protocol.
25   *
26   * @param msg Batch of Transactions of type ClientQueryBatch from the client.
27   * @return RC
28   */
29 RC WorkerThread::process_client_batch(Message *msg)
30 {
31     //printf("ClientQueryBatch: %ld, THD: %ld :: CL: %ld :: RQ: %ld\n", msg->txns->size(), msg->txns->begin() - msg->txns, msg->txns->begin(), msg->txns->begin());
32     //fflush(stdout);
33
34     ClientQueryBatch *clbtch = (ClientQueryBatch *)msg;
35
36     // Authenticate the client signature.
37     validate_msg(clbtch);
38
39 #if VIEW_CHANGES
40     // If message forwarded to the non-primary.
41     if (g_node_id != get_current_view(get_thd_id()))
42     {
43         client_query_check(clbtch);
44         return RCOKE;
45     }
46
47     // Partial failure of Primary 0.
48     fail_primary(msg, 9);
49 #endif
50
51     // Initialize all transaction managers and uint64_t Message::txnid.
52     create_and_send_batchreq(clbtch, clbtch->txnid);
53
54     return RCOKE;
55 }
```

```
C++ worker_thread.cpp ×
system > C++ worker_thread.cpp > WorkerThread::create_and_send_batchreq(ClientQueryBatch *, uint64_t)
1123  * This function is used by the primary replicas to create and set
1124  * transaction managers for each transaction part of the ClientQueryBatch message
1125  * by the client. Further, to ensure integrity a hash of the complete batch is
1126  * generated, which is also used in future communication.
1127  *
1128  * @param msg Batch of transactions as a ClientQueryBatch message.
1129  * @param tid Identifier for the first transaction of the batch.
1130  */
1131 void WorkerThread::create_and_send_batchreq(ClientQueryBatch *msg, uint64_t tid)
1132 {
1133     // Creating a new BatchRequests Message.
1134     Message *bmsg = Message::create_message(BATCH_REQ);
1135     BatchRequests *breq = (BatchRequests *)bmsg;
1136     breq->init(get_thd_id());
1137
1138     // Starting index for this batch of transactions.
1139     next_set = tid;
1140
1141     // String of transactions in a batch to generate hash.
1142     string batchStr;
1143
1144     // Allocate transaction manager for all the requests in batch.
1145     for (uint64_t i = 0; i < get_batch_size(); i++)
1146     {
1147         uint64_t txn_id = get_next_txnid() + i;
1148
1149         //cout << "Txn: " << txn_id << " :: Thd: " << get_thd_id() << "\n";
1150         //fflush(stdout);
1151         txn_man = get_transaction_manager(txn_id, 0);
1152
1153         // Unset this txn man so that no other thread can concurrently use.
1154         while (true)
1155         {
1156             bool ready = txn_man->unset_ready();
1157             if (!ready)
1158             {
1159                 continue;
1160             }
1161             else
1162             {
1163                 break;
1164             }
1165         }
1166     }
1167 }
```

PBFT: Practical Byzantine Fault Tolerance

Process Batch Request (Prepare)

- System/worker_thread_pbft.cpp
- process_batch Function
- Create and Send Prepare Message
 - Create Transactions
 - Save Digest
- PBFTPrepare Class
 - Prepare Message

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > WorkerThread::process_batch(Message *)
57  /**
58   * Process incoming BatchRequests message from the Primary.
59   *
60   * This function is used by the non-primary or backup replicas to process an incoming
61   * BatchRequests message sent by the primary replica. This processing would require
62   * sending messages of type PBFTPrepMessage, which correspond to the Prepare phase of
63   * the PBFT protocol. Due to network delays, it is possible that a replica may have
64   * received some messages of type PBFTPrepMessage and PBFTCommitMessage, prior to
65   * receiving this BatchRequests message.
66   *
67   * @param msg Batch of Transactions of type BatchRequests from the primary.
68   * @return RC
69   */
70 RC WorkerThread::process_batch(Message *msg)
71 {
72     uint64_t cntime = get_sys_clock();
73
74     BatchRequests *breq = (BatchRequests *)msg;
75
76     //printf("BatchRequests: TID:%ld : VIEW: %ld : THD: %ld\n", breq->txn_id, breq->view, get_
77     //fflush(stdout);
78
79     // Assert that only a non-primary replica has received this message.
80     assert(g_node_id != get_current_view(get_thd_id()));
81
82     // Check if the message is valid.
83     validate_msg(breq);
84 }
```

PBFT: Practical Byzantine Fault Tolerance

Process Prepare and Commit Messages(Prepare)

- System/worker_thread_pbft.cpp
- process_pbft_prepare Function
 - Count Prepare Messages
 - Create and Send commit Message
 - PBFTCommit Message
- process_pbft_commit Function
 - Count commit messages
 - Create and Send execute Message
 - ExecuteMessage Class

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > ...
200
186 /**
187 * Processes incoming Prepare message.
188 *
189 * This function processes incoming messages of type PBFTPrepMessage. If
190 * received 2f identical Prepare messages from distinct replicas, then it c
191 * and sends a PBFTCommitMessage to all the other replicas.
192 *
193 * @param msg Prepare message of type PBFTPrepMessage from a replica.
194 * @return RC
195 */
196 RC WorkerThread::process_pbft_prep_msg(Message *msg)
197 {
198     //cout << "PBFTPrepMessage: TID: " << msg->txnid << " FROM: " << msg->
199     //fflush(stdout);
200
201     // Start the counter for prepare phase.
202     if (txn_man->prep_rsp_cnt == 2 * g_min_invalid_nodes)
203     {
204         txn_man->txn_stats.time_start_prepare = get_sys_clock();
205     }
206
207     // Check if the incoming message is valid.
208     PBFTPrepMessage *pmsg = (PBFTPrepMessage *)msg;
209     validate_msg(pmsg);
210
211     // Check if sufficient number of Prepare messages have arrived.
212     if (prepared(pmsg))
213     {
214         // Send Commit messages.
215         txn_man->send_pbft_commit_msgs();
216
217         // End the prepare counter.
218         INC_STATS(get_thd_id(), time_prepare, get_sys_clock() - txn_man->tx
219     }
220
221     return RCOK;
222 }
```

```
C++ worker_thread_pbft.cpp ×
system > C++ worker_thread_pbft.cpp > ⚡ WorkerThread::process_pbft_commit_msg(Message *)
200
275 /**
276 * Processes incoming Commit message.
277 *
278 * This function processes incoming messages of type PBFTCommitMessage
279 * received 2f+1 identical Commit messages from distinct replicas, then :
280 * execute-thread to execute all the transactions in this batch.
281 *
282 * @param msg Commit message of type PBFTCommitMessage from a replica.
283 * @return RC
284 */
285 RC WorkerThread::process_pbft_commit_msg(Message *msg)
286 {
287     //cout << "PBFTCommitMessage: TID: " << msg->txnid << " FROM: " << msg->
288     //fflush(stdout);
289
290     if (txn_man->commit_rsp_cnt == 2 * g_min_invalid_nodes + 1)
291     {
292         txn_man->txn_stats.time_start_commit = get_sys_clock();
293     }
294
295     // Check if message is valid.
296     PBFTCommitMessage *pcmsg = (PBFTCommitMessage *)msg;
297     validate_msg(pcmsg);
298
299     txn_man->add_commit_msg(pcmsg);
300
301     // Check if sufficient number of Commit messages have arrived.
302     if (committed_local(pcmsg))
303     {
304 #if TIMER_ON
305         // End the timer for this client batch.
306         server_timer->endTimer(txn_man->hash);
307    #endif
308
309     // Add this message to execute thread's queue.
310     send_execute_msg();
311
312     INC_STATS(get_thd_id(), time_commit, get_sys_clock() - txn_man->tx
313 }
```

PBFT: Practical Byzantine Fault Tolerance

Process Execute Message

- System/worker_thread.cpp
- Internal Message
- process_execute Function
- Execute the Transactions in batch in order
- Create and send Client Response
- ClientResponse Class

```
C++ worker_thread.cpp X
system > C++ worker_thread.cpp > WorkerThread::process_execute_msg(Message *)
795
796 /**
797 * Execute transactions and send client response.
798 *
799 * This function is only accessed by the execute-thread, which executes the transactions
800 * in a batch, in order. Note that the execute-thread has several queues, and at any
801 * point of time, the execute-thread is aware of which is the next transaction to
802 * execute. Hence, it only loops on one specific queue.
803 *
804 * @param msg Execute message that notifies execution of a batch.
805 * @ret RC
806 */
807 RC WorkerThread::process_execute_msg(Message *msg)
808 {
809     //cout << "EXECUTE " << msg->txnid << " :: " << get_thd_id() <<"\n";
810     //fflush(stdout);
811
812     uint64_t ctime = get_sys_clock();
813
814     // This message uses txnid of index calling process_execute.
815     Message *rsp = Message::create_message(CL_RSP);
816     ClientResponseMessage *crsp = (ClientResponseMessage *)rsp;
817     crsp->init();
818
819     ExecuteMessage *emsg = (ExecuteMessage *)msg;
820
821     // Execute transactions in a shot
822     uint64_t i;
823     for (i = emsg->index; i < emsg->end_index - 4; i++)
824     {
825         //cout << "i: " << i << " :: next index: " << g_next_index << "\n";
826         //fflush(stdout);
827
828         TxnManager *tman = get_transaction_manager(i, 0);
829
830         inc_next_index();
831
832         // Execute the transaction
833         tman->run_txnid();
834
835         // Commit the results.
836         tman->commit();
837
838         crsp->copy_from_txn(tman);
```

PBFT: Practical Byzantine Fault Tolerance

Work Queue

- Lock Free queues
- All the messages are being stored in these queues
- System/work_queue.cpp
- Multiple queues for different Threads
- Dequeue and Enqueue Interfaces
- Enqueue in IOThread
- Dequeue in Worker Thread

```
C++ work_queue.cpp ×  
system > C++ work_queue.cpp > ...  
44     void QWorkQueue::enqueue(uint64_t thd_id, Message *msg, bool busy)  
45     {  
46         uint64_t starttime = get_sys_clock();  
47         assert(msg);  
48         DEBUG_M("QWorkQueue::enqueue work_queue_entry alloc\n");  
49         work_queue_entry *entry = (work_queue_entry *)mem_allocator.align_alloc(sizeof(work_queue_ent  
50             entry->msg = msg;  
51             entry->rtype = msg->rtype;  
52             entry->txn_id = msg->txn_id;  
53             entry->batch_id = msg->batch_id;  
54             entry->starttime = get_sys_clock();  
55             assert(ISSERVER || ISREPLICA);  
56             DEBUG("Work Enqueue (%ld,%ld) %d\n", entry->txn_id, entry->batch_id, entry->rtype);  
57  
58             if (msg->rtype == CL_QRY || msg->rtype == CL_BATCH)  
59             {  
60                 if (g_node_id == get_current_view(thd_id))  
61                 {  
62                     //cout << "Placing \n";  
63                     while (!new_txn_queue->push(entry) && !simulation->is_done())  
64                     {  
65                         //cout << "Pushing \n";  
66                     }  
67                 else  
68                 {  
69                     assert(entry->rtype < 100);  
70                     while (!work_queue[0]->push(entry) && !simulation->is_done())  
71                     {  
72                         //cout << "Pushing \n";  
73                     }  
74                 }  
75             }  
76         }  
77     }  
78 }
```

PBFT: Practical Byzantine Fault Tolerance

IO Thread and Transport Layer

- Multiple Input Threads
- Multiple Output Threads
- System/io_thread.cpp
- Transport Layer: TCP Sockets
- Nano Message Library
- Transport/transport.cpp

```
C++ io_thread.cpp ×  
system > C++ io_thread.cpp > ...  
299     RC InputThread::server_recv_loop()  
300     {  
301  
302         myrand rdm;  
303         rdm.init(get_thd_id());  
304         RC rc = RCOK;  
305         assert(rc == RCOK);  
306         uint64_t starttime = 0;  
307         uint64_t idle_starttime = 0;  
308         std::vector<Message *> *msgs;  
309         while (!simulation->is_done())  
310         {  
311             heartbeat();  
312  
313             #if VIEW_CHANGES  
314                 if (g_node_id != get_current_view(get_thd_id()))  
315                 {  
316                     uint64_t tid = get_thd_id() - 1;  
317                     uint32_t nchange = get_newView(tid);  
318  
319                     if (nchange)  
320                     {  
321                         set_current_view(get_thd_id(), get_current_view(get_thd_id()) + 1);  
322                         set_newView(tid, false);  
323                     }  
324                 }  
325             #endif  
326  
327             msgs = tport_man.recv_msg(get_thd_id());  
328         }
```

Configuration Parameters to Play

- NODE_CNT
Total number of replicas, minimum 4, that is, $f=1$.
- THREAD_CNT
Total number of threads at primary (at least 5)
- CLIENT_NODE_CNT
Total number of clients (at least 1).
- MAX_TXN_IN_FLIGHT
Multiple of Batch Size
- DONE_TIMER
Amount of time to run the system.
- BATCH_THREADS
Number of threads at primary to batch client transactions.
- BATCH_SIZE
Number of transactions in a batch (at least 10)
- TXN_PER_CHKPT
Frequency at which garbage collection is done.
- USE_CRYPTO
To switch on and off cryptographic signing of messages.
- CRYPTO_METHOD_ED25519
To use ED25519 based digital signatures.
- CRYPTO_METHOD_CMAC_AES
To use CMAC + AES combination for authentication

Thank You